

Newton Extensions

Minghe

Quasi-Newton Methods (Surrogate Hessian Approximations)

SR1 (Symmetric Rank-1 Update)

```
quasiNewton_SR1 <- function(x0, grad, tol=1e-10, max_iter=500) {  
  # x0: initial parameter vector (of length p)  
  p <- length(x0)  
  x <- x0  
  H_inv <- diag(p) # initial inverse Hessian approximation (identity)  
  g <- grad(x)  
  i <- 0  
  res <- list(x = x, iter = i, grad = g)  
  while(i < max_iter && sqrt(sum(g^2)) > tol) {  
    s <- - H_inv %*% g # proposed step  
    x_new <- x + s  
    g_new <- grad(x_new)  
    y_vec <- g_new - g  
    d <- s - H_inv %*% y_vec  
    denom <- as.numeric(t(d) %*% y_vec)  
    if(abs(denom) > 1e-8) {  
      H_inv <- H_inv + (d %*% t(d)) / denom  
    }  
    x <- x_new  
    g <- g_new  
    i <- i + 1  
    res[[i+1]] <- list(x = x, iter = i, grad = g)  
  }  
  return(res)  
}
```

DFP (Davidon-Fletcher-Powell)

```
quasiNewton_DFP <- function(x0, grad, tol=1e-10, max_iter=500) {  
  p <- length(x0)  
  x <- x0  
  H_inv <- diag(p)  
  g <- grad(x)  
  i <- 0  
  res <- list(x = x, iter = i, grad = g)  
  while(i < max_iter && sqrt(sum(g^2)) > tol) {  
    s <- - H_inv %*% g  
    x_new <- x + s  
    g_new <- grad(x_new)
```

```

y_vec <- g_new - g
s <- matrix(s, ncol=1)
y_vec <- matrix(y_vec, ncol=1)
H_inv <- H_inv + (s %*% t(s))/(t(s) %*% y_vec) - (H_inv %*% y_vec %*% t(y_vec) %*% H_inv)/(t(y_vec) %*% y_vec)
x <- x_new
g <- g_new
i <- i + 1
res[[i+1]] <- list(x = x, iter = i, grad = g)
}
return(res)
}

```

BFGS (Broyden–Fletcher–Goldfarb–Shanno)

```

quasiNewton_BFGS <- function(x0, grad, tol=1e-10, max_iter=500) {
  p <- length(x0)
  x <- x0
  H_inv <- diag(p)
  g <- grad(x)
  i <- 0
  res <- list(x = x, iter = i, grad = g)
  while(i < max_iter && sqrt(sum(g^2)) > tol) {
    s <- - H_inv %*% g
    x_new <- x + s
    g_new <- grad(x_new)
    y_vec <- g_new - g
    rho <- as.numeric(1 / (t(y_vec) %*% s))
    if(rho <= 0) {
      # Reset if update is not positive-definite
      H_inv <- diag(p)
    } else {
      I <- diag(p)
      H_inv <- (I - rho * s %*% t(y_vec)) %*% H_inv %*% (I - rho * y_vec %*% t(s)) + rho * s %*% t(s)
    }
    x <- x_new
    g <- g_new
    i <- i + 1
    res[[i+1]] <- list(x = x, iter = i, grad = g)
  }
  return(res)
}

```

Coordinate–Wise Optimization

```

coordinateDescent <- function(x0, f, grad, tol=1e-10, max_iter=500) {
  x <- x0
  p <- length(x)
  i <- 0
  while(i < max_iter) {
    x_old <- x
    for(j in 1:p) {
      # Here, we update the jth coordinate while holding others fixed.

```

```

    # One approach is to compute the partial derivative w.r.t. x[j] and take a small step.
    g <- grad(x)
    step <- 1e-3 # A fixed step size (or use a line search)
    x[j] <- x[j] - step * g[j]
  }
  i <- i + 1
  if(sqrt(sum((x - x_old)^2)) < tol) break
}
return(list(x = x, iter = i))
}

```

path-wise coordinate-descent algorithm for the logistic-lasso

```

# Soft-thresholding operator (for L1 penalty)
soft_threshold <- function(z, gamma) {
  sign(z) * pmax(abs(z) - gamma, 0)
}

# The main function to compute the logistic-lasso path using coordinate descent
logistic_lasso_path <- function(X, y, lambda_min_ratio = 0.001, nlambda = 100,
                                tol = 1e-4, max_iter = 1000) {
  # X is an n x p matrix (predictors) and y is a binary vector (0/1)
  n <- nrow(X)
  p <- ncol(X)

  # Step 1: Compute lambda_max. For logistic regression with intercept,
# when beta = 0 we have p_i = 1/2, so the subgradient condition yields:
# |sum_i x_ij (y_i - 1/2)| <= lambda, for all j.
  lambda_max <- max(abs(colSums(X * (y - 0.5))))

  # Step 2: Define a decreasing sequence of lambda values from lambda_max to lambda_min.
  lambda_min <- lambda_max * lambda_min_ratio
  lambda_seq <- exp(seq(log(lambda_max), log(lambda_min), length.out = nlambda))

  # Initialize coefficient estimates.
# For the intercept, a common starting value is the logit of the mean of y.
  beta0 <- log(mean(y) / (1 - mean(y)))
  beta <- rep(0, p)

  # To store the path: one row per lambda.
  intercept_path <- numeric(nlambda)
  beta_path <- matrix(0, nrow = nlambda, ncol = p)

  # Loop over lambda values
  for (l in 1:nlambda) {
    lambda <- lambda_seq[l]

    # Coordinate descent for current lambda.
    converged <- FALSE
    iter <- 0
    while (!converged && iter < max_iter) {
      iter <- iter + 1
    }
  }
}

```

```

# Compute current linear predictor and probabilities.
eta <- beta0 + X %*% beta
p_i <- 1 / (1 + exp(-eta))

# Compute weights and working response.
w <- p_i * (1 - p_i)
# To avoid division by very small w, one might add a small constant if needed.
z <- eta + (y - p_i) / w

# Update the intercept (unpenalized):
beta0_new <- sum(w * (z - X %*% beta)) / sum(w)

# Update each coordinate for beta:
beta_new <- beta
for (j in 1:p) {
  # Compute partial residual excluding feature j:
  r_j <- z - beta0_new - X[, -j, drop = FALSE] %*% beta_new[-j]
  # Compute numerator and denominator for the j-th coordinate.
  num <- sum(w * X[, j] * r_j)
  den <- sum(w * X[, j]^2)
  # Apply soft-thresholding update:
  beta_new[j] <- soft_threshold(num, lambda) / den
}

# Check for convergence (e.g. maximum change in coefficients is below tol)
if (max(abs(beta0_new - beta0), abs(beta_new - beta)) < tol) {
  converged <- TRUE
}

beta0 <- beta0_new
beta <- beta_new
} # end of coordinate descent for current lambda

intercept_path[l] <- beta0
beta_path[l, ] <- beta
}

return(list(lambda_seq = lambda_seq, intercept_path = intercept_path,
            beta_path = beta_path))
}

# Suppose X is an n x p matrix and y is an n-vector (with values 0 or 1)
# fit_path <- logistic_lasso_path(X, y, lambda_min_ratio = 0.001, nlambda = 100)
# You can inspect fit_path$lambda_seq, fit_path$intercept_path, and fit_path$beta_path

```